

The Eye

A data ingestion pipeline for casino gaming analytics

W205 Final Presentation

Mohammad J. Habib



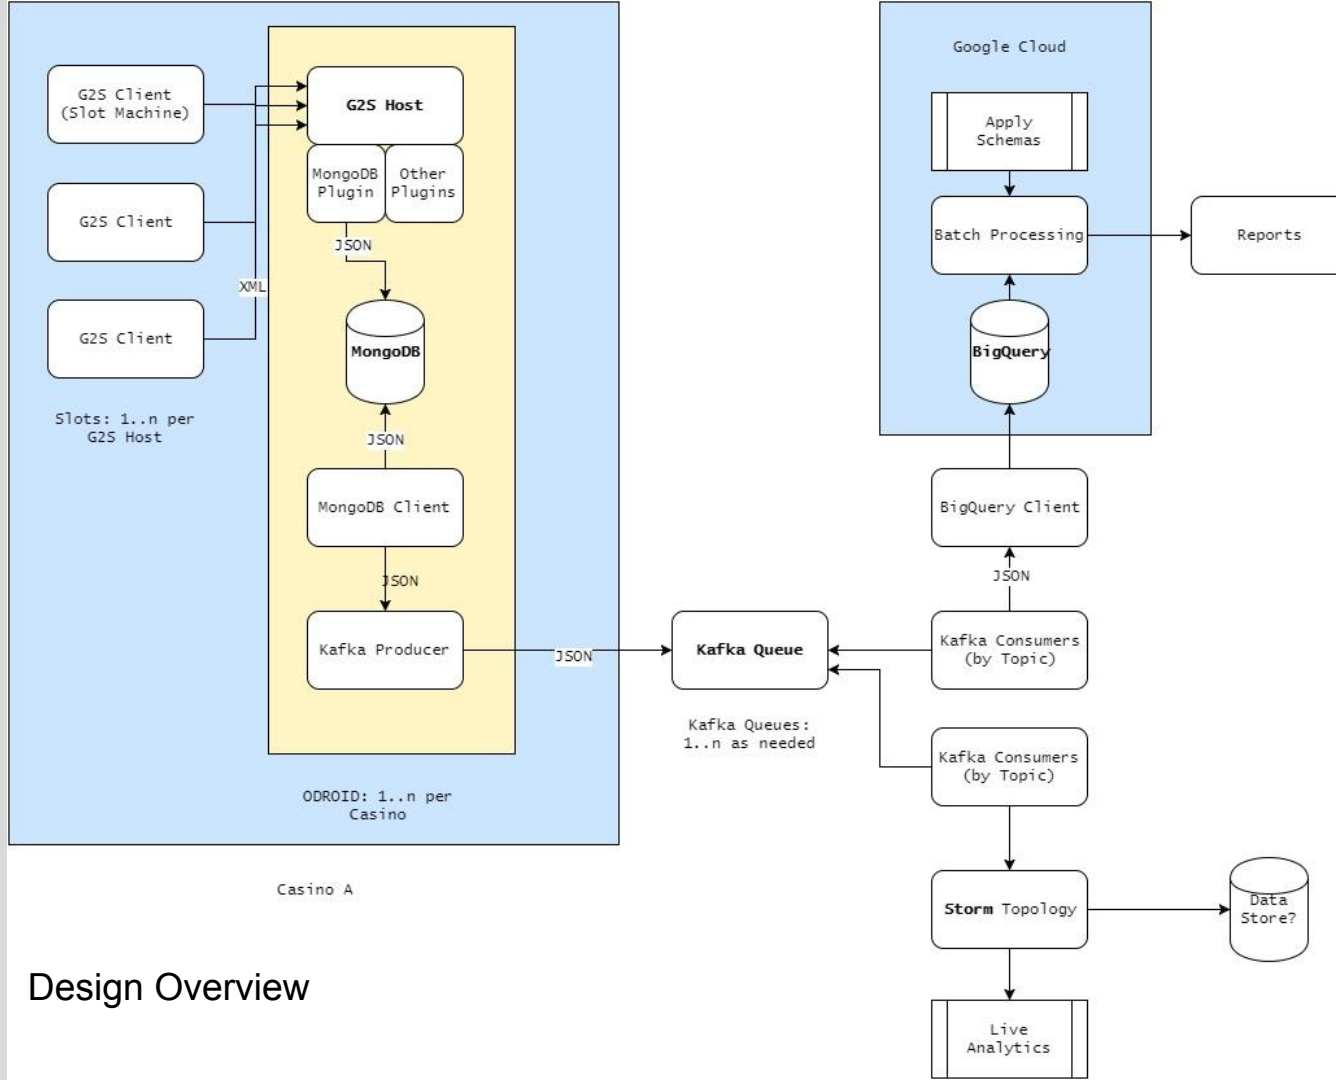
Problem Overview

Getting data from slot machines is very difficult for vendors:

- Regulators strictly control access to gaming machines
- Casino operators guard accounting data gathered from slot machines
- Casinos traditionally have Serial networks (not Ethernet); Web connectivity is spotty at best
- Using data for billing and game performance analytics is not easily possible

Plan for getting data

- Introduce a cheap hardware device that can “attach” to EGMs and listen for data (ODROID C1+)
- Use G2S; a common protocol used by many slot machine vendors
- Use VPN over cellular network to get data out of Casinos
- Batch data locally and send when connection becomes available



Scalability

Architecture scales horizontally:

- Add ODROIDS to support more Slot machines
- Add more Kafka queues to ingest higher volume of data
- Add more consumers as range of topics and volume grows
- BigQuery datasets can be created per Casino or Region
- BigQuery tables can be partitioned by date
- Add storage or remove old data as needed

Extensibility

- Stream processing can be added for live analytics (just add a new consumer)
- Other data processing systems can be bolted on e.g. Spark
- Bigquery data can be exposed to Tableau, Google Sheets and other business apps
- Bigquery can be replaced with another backend e.g. Hadoop

Known Limitations

The “proof of concept” developed for this project has following limitations. I’ve also listed how these limitations can be addressed for production.

Limitation in Proof of Concept	How to address in Production?
Data are not assigned a unique ODROID identifier and therefore we cannot tell apart data from one ODROID and another	<ul style="list-style-type: none">• Add an ODROID identifier to each row of data e.g. ODROID public IP or MAC address, Or• Store data from each different ODROID in a different table
Data are currently not removed from the MongoDB on ODROID	<ul style="list-style-type: none">• Purge data after a certain number of days from MongoDB to free up disk space on ODROID
Data is batched up in an in-memory Python queue before sending to BigQuery	<ul style="list-style-type: none">• Replace in-memory queue with something like Redis
BigQuery tables are not partitioned (one table is forever growing)	<ul style="list-style-type: none">• Partition tables based on date, or by ODROID + Date
Storm topology is not yet implemented (lack of time)	<ul style="list-style-type: none">• This can be implemented in the future for production

Complexity and Storage

Architecture is fairly straightforward and decoupled:

- Data comes in as XML, gets transformed to JSON and makes its way through the pipeline. It is stored as JSON in BigQuery
- Schemas are applied in BigQuery and the data can be combined with additional data (e.g. marketing identifiers for games)
- Complexity can be further reduced by RESTifying the data exchange between various components

Storage and throughput requirements are fairly demanding:

- Load test showed that 10 Slots generate ~300MB of data in 12 hours when set to auto-play (one wager per 3 seconds per slot).
- Guessing that we will need to move ~150MB per Casino per day. Since Slots are not constantly in auto-play.
- There are 1000 or so Casinos who are Aristocrat's customers. Overall ~150GB of data moved and stored per day.

Google BigQuery



New Query ?

Query Editor UDF Editor ✕

```
1 SELECT
2   egmId, wagerCategory,
3   COUNT(wagerCategory) AS wagerCategoryCounts
4 FROM (
5   SELECT
6     _id,
7     REGEXP_EXTRACT(Key, r'^([~]*)') AS egmId,
8     JSON_EXTRACT(Value, '$.meterList.meterInfo[1].wagerMeters[0].wagerCategory') AS wagerCategory
9 FROM (
10  SELECT
11    *
```

RUN QUERY Save Query Save View Format Query Show Options Query complete (0.8s elapsed, cached)

Results Explanation Download as CSV Download as JSON Save as Table Save to Google Sheets

Row	egmId	dateTime	wager	win	
1	ATI_GEN7_3064128B17	2016-07-21 15:53:06 UTC	1.25	5.0	
2	ATI_GEN7_3064128B17	2016-07-21 15:53:20 UTC	1.25	1.75	
3	ATI_GEN7_3064128B17	2016-07-21 15:53:24 UTC	1.25	0.0	
4	ATI_GEN7_3064128B17	2016-07-21 15:53:28 UTC	1.25	0.0	
5	ATI_GEN7_3064128B17	2016-07-21 15:53:31 UTC	1.25	0.0	
6	ATI_GEN7_3064128B17	2016-07-21 15:53:33 UTC	1.25	0.0	
7	ATI_GEN7_3064128B17	2016-07-21 15:55:01 UTC	1.25	32.75	
8	ATI_GEN7_3064128B17	2016-07-21 15:59:20 UTC	1.25	0.0	
9	ATI_GEN7_3064128B17	2016-07-21 16:04:59 UTC	1.25	0.0	
10	ATI_GEN7_3064128B17	2016-07-21 16:05:06 UTC	1.25	4.0	
11	ATI_GEN7_3064128B17	2016-07-21 16:05:07 UTC	1.25	0.0	

Google BigQuery



New Query ?

Query Editor UDF Editor ✕

```
1 SELECT * FROM [sauronbigquery:sauron.view_winLossByEgm] LIMIT 1000
```

Valid: This query will process 422 MB when run.

RUN QUERY ▼

Save Query

Save View

Format Query

Show Options

Query complete (3.3s elapsed, 422 MB processed)



Results Explanation

Download as CSV

Download as JSON

Save as Table

Save to Google Sheets

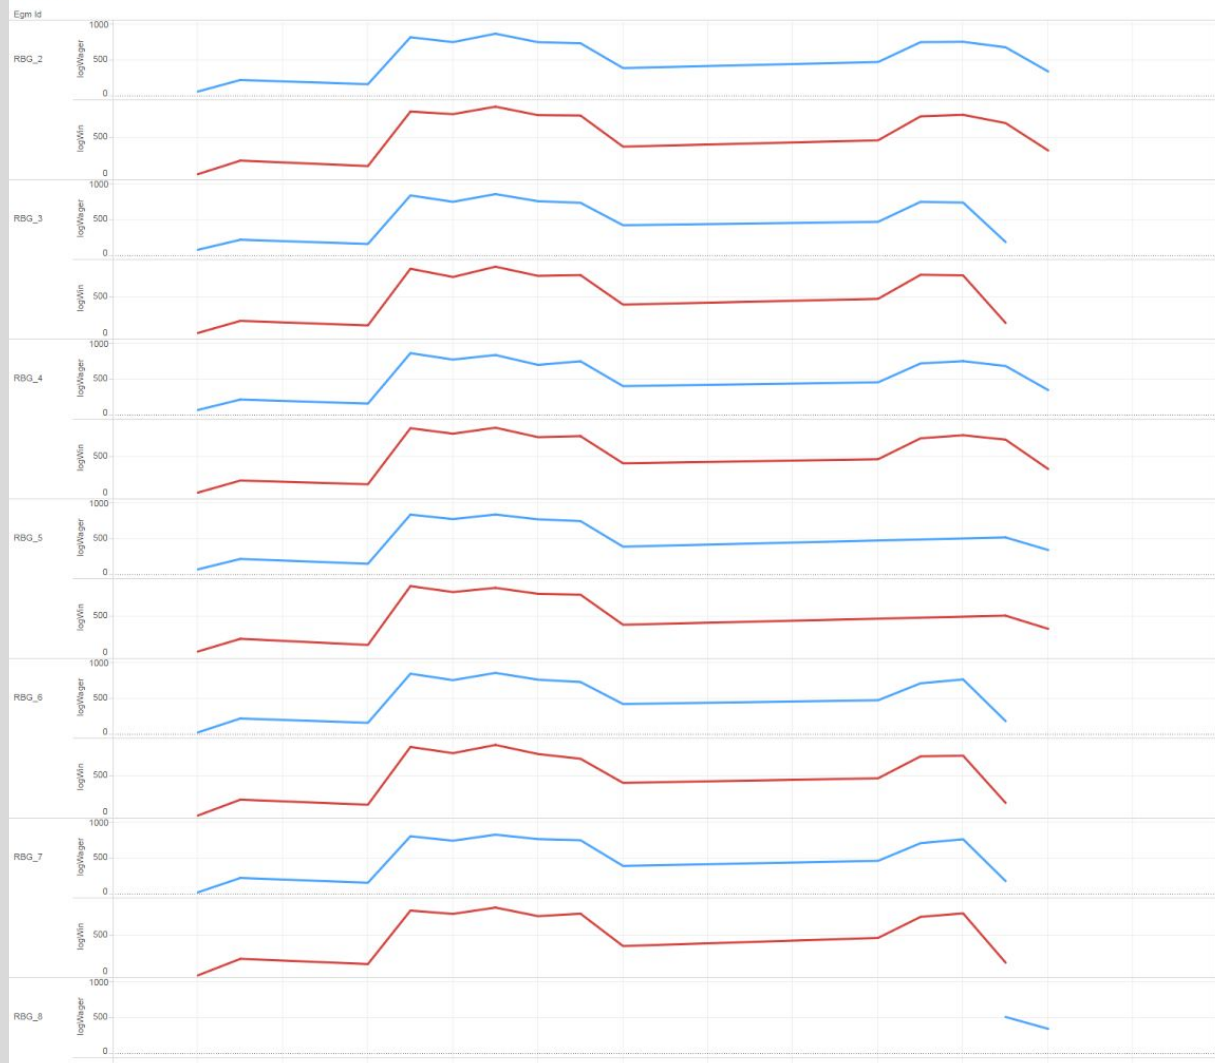
Row	egmId	totalBet	totalWin	netWin	winRatio	
1	ATI_GEN7_3064128B17	885625000	878725000	-6900000	0.9922088920253944	
2	RBG_1234	1026600000	53870000	-972730000	0.05247418663549903	

Table JSON


wagersWinsHistogram



logWinsWagersPerHour



netPlayerWinPerHour



The trend of sum of netPlayerLoss for Date Time Hour broken down by Egm Id. Color shows details about Egm Id.

cumulativePlayerWinOverTime

